

# CS 2200 Spring 2011 Test 2

Prism ID: \_\_\_\_\_

Name: \_\_\_\_\_ GTID#: 9 \_\_\_\_\_

Problem	Points	Lost	Gained	Running Total	TA
1	1				
2	15				
3	5				
4	15				
5	14				
6	20				
7	5				
8	10				
9	15				
Total	100				

- You may ask for clarification but you are ultimately responsible for the answer you write on the paper.
- Illegible answers are wrong answers.
- Please do not discuss this test by any means (until 5 pm today)
- Please look through the entire test before starting. WE MEAN IT!!!

**Illegible answers are wrong answers.**

Good luck!

1. (1 point, 1 min) (circle one)

Which of the following teams are in the semifinals of the **2011 World Cup Cricket** (circle all that apply)?

- |                |                 |                |
|----------------|-----------------|----------------|
| a) Australia   | b) South Africa | c) India       |
| d) New Zealand | e) Pakistan     | f) USA         |
| g) Sri Lanka   | h) England      | i) West Indies |

# CS 2200 Spring 2011 Test 2

Prism ID: \_\_\_\_\_

Name: \_\_\_\_\_ GTID#: 9 \_\_\_\_\_

## Pipelining

2. (15 points, 10 min)

(a) (2 points) (Select one **correct** choice)

With reference to the 5-stage pipelined implementation of LC-2200 instruction set, a second ALU is needed in the EXEC stage of the pipeline

- 1) To make the stage symmetrical
- 2) To enable address calculation for LW/SW instructions
- 3) To enable address calculation for BEQ instruction
- 4) To enable multi-precision arithmetic for ADD instruction
- 5) To enable ADD and NAND instructions to be executed in parallel if they occur one after another

(b) (3 points)

What role does Branch Target Buffer play in a pipelined processor design?

(c) (5 points)

**I<sub>1</sub>: R1 ← R2 + R3**

**I<sub>2</sub>: R4 ← R1 + R5**

If I<sub>2</sub> is immediately following I<sub>1</sub> in the pipeline with **no forwarding**, how many bubbles will be experienced by the pipeline? **You must explain your answer to get any credit.**

# CS 2200 Spring 2011 Test 2

Name: \_\_\_\_\_ Prism ID: \_\_\_\_\_  
GTID#: 9 \_\_\_\_\_

(d) (5 points)

**I<sub>1</sub>: LW R1, 0(R2); load R1 from memory**

**I<sub>2</sub>: R4 <- R1 + R5**

If I<sub>2</sub> is immediately following I<sub>1</sub> in the pipeline with **data forwarding**, how many bubbles will be experienced by the pipeline? **You must explain your answer to get any credit.**

## Process Scheduling

3. (5 points, 5 min)

(Use each of the following terms EXACTLY ONCE: "FCFS", "waiting", "SJF", "ready", "PCB", "disk image", "round robin". NOTE: NOT ALL TERMS ARE USED.)

- a) If process A is currently running and makes a request to read a large block of data from the disk, the OS puts this process in the \_\_\_\_\_ state.
- b) If process A is currently running and the OS preempts A since its time quantum is up to run a different process B, then the OS puts process A in the \_\_\_\_\_ state.
- c) To dispatch a process A to run on the processor, the OS must take register values from A's \_\_\_\_\_ and load them into actual processor registers.
- d) If fairness is more important than average waiting time experienced by the processes, we should use \_\_\_\_\_ non-preemptive scheduling algorithm.
- e) If we want to optimize the average waiting time for processes, then we should use \_\_\_\_\_ non-preemptive scheduling algorithm.

# CS 2200 Spring 2011 Test 2

Prism ID: \_\_\_\_\_

Name: \_\_\_\_\_ GTID#: 9 \_\_\_\_\_

4. (15 points, 15 min)

Recall that *Shortest Remaining Time First (SRTF)* is a variant of Shortest Job First (SJF) with preemption added in. Consider the following three processes vying for the CPU. The scheduler uses SRTF. The scheduler re-evaluates which process to run only upon the arrival of a new process into the scheduling queue, or the completion of a process. The table shows the arrival time of each process.

Process	Arrival Time	Execution Time
P1	$T_0$	2ms
P2	$T_0+1\text{ms}$	3ms
P3	$T_0+2\text{ms}$	1ms

The scheduling starts at time  $T_0$ .

- (a) (6 points) Fill in the table below with the process that is executing on the processor during each time slot.

Interval $T_0+$	0	1	2	3	4	5	6	7	8	9	10	11	12
Running													

Use the following tables to show your work as to how you arrived at the above schedule.

Time  $T_0$ :

Process	Remaining time
P1	2ms
P2	Not arrived yet
P3	Not arrived yet

Time  $T_0+1$ : (2 points) (Process P2 arrives)

Process	Remaining time
P1	
P2	
P3	Not arrived yet

Time  $T_0+2$ : (3 points) (Process P3 arrives)

Process	Remaining time
P1	
P2	
P3	

# CS 2200 Spring 2011 Test 2

Name: \_\_\_\_\_ Prism ID: \_\_\_\_\_  
GTID#: 9 \_\_\_\_\_

(b) (2 points) What is the total waiting time for all processes?

(c) (2 points) What will be the total waiting time for all processes with an FCFS schedule?

## Memory Management and Virtual Memory

5. (14 points, 10 min)

(a) (2 points) (choose one of the following)

In a byte-addressed machine, if the pagesize is 8192 bytes, the **maximum** possible internal fragmentation is

- 1) 8192 bytes
- 2) 1 byte
- 3) 8191 bytes
- 4) There is no internal fragmentation with paging

(b) (4 points)

Virtual address is 32 bits; pagesize 8192 bytes (8 Kbytes); How many entries are there in the page table?

(c) (4 points)

For the same memory system as in (b), the physical address is 24 bits. How many physical page frames does the memory system have?

# CS 2200 Spring 2011 Test 2

Prism ID: \_\_\_\_\_

Name: \_\_\_\_\_ GTID#: 9 \_\_\_\_\_

(d) (2 points)

Total memory size is 24K. The current state of the memory allocation table for a variable-sized partition allocation scheme is as follows:

Start address	Size	Who has it?
0	4K	P1
4096	4K	P2
8192	4K	FREE
12288	2K	P3
14336	10K	FREE

What is the total amount of external fragmentation?

(e) (2 points)

The necessary conditions for a paging memory system are (select one of the following choices)

- 1) The virtual and physical addresses have to be of the same size
- 2) The sizes of the virtual page and the physical frame have to be the same
- 3) The number of physical frames should be larger than the number of virtual pages
- 4) (1) and (2)
- 5) (2) and (3)
- 6) (1) and (3)
- 7) (1), (2), and (3)

## Working set, page replacement, TLB

6. (20 points, 15 min)

During the time interval  $t_1 - t_2$ , the following virtual page accesses are recorded for the process P1.

P1: 0, 10, 1, 0, 1, 2, 10, 2, 1, 1, 0

a) (2 points) What is the **working set** for P1 in this time interval?

b) (3 points) What is the **memory pressure** on the system during this interval?

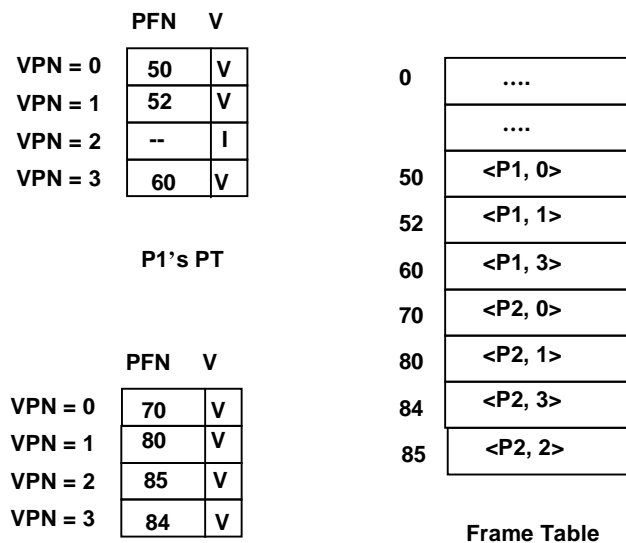
# CS 2200 Spring 2011 Test 2

Prism ID: \_\_\_\_\_

Name: \_\_\_\_\_ GTID#: 9 \_\_\_\_\_

c) (6 points)


P1 incurs a page fault at VPN = 2. The victim page frame chosen by the page replacement algorithm is PFN = 84. The "before" picture of the page manager's data structure are shown below. Note that only relevant entries of the Frame Table are shown in the Figures below. Show the contents of the data structures "after" the page fault is serviced. **You can simply mark the changes in the figure below to show your answer.**



d) (9 points)

A processor has a 4-entry TLB that uses a least recently used (LRU) replacement policy. The program executed by this processor generates the following sequence of memory requests (only the page number of each request is shown). Assuming that the TLB is initially empty, which of these memory requests will result in TLB misses?

LRU stack



# CS 2200 Spring 2011 Test 2

Prism ID: \_\_\_\_\_

Name: \_\_\_\_\_ GTID#: 9 \_\_\_\_\_

Page #	TLB Miss?
0xF000	
0xE000	
0x10AB	
0xF000	
0x2000	
0x2001	
0x2000	
0xF000	
0x4000	
0xE000	
0x10AB	
0x2001	
0x10AB	
0x2000	
0x2001	

## Caching

7. (5 points, 5 min)

A pipelined processor has an average CPI of 1.2 not considering memory effects. On an average each instruction has an I-cache miss of 1%, and a D-cache miss of 1%. The miss penalty is 20 cycles. What is the effective CPI taking into account memory stalls?

8. (10 points, 10 min)

Consider the following memory hierarchy:

L1 cache: Access time = 2ns; hit rate = 99%

L2 cache: Access time = 10ns; hit rate = 95%

Main memory: Access time = 100ns

Compute the effective memory access time.



# CS 2200 Spring 2011 Test 2

Name: \_\_\_\_\_ Prism ID: \_\_\_\_\_  
GTID#: 9 \_\_\_\_\_

9. (15 points, 15 min)

Consider a 2-way set-associative cache for byte addressed processor.

- Data size of cache = 16KB.
- CPU address = 32 bits
- Memory word = 4 bytes.
- Cache block size = 16 bytes.
- Write policy is Write-back at the granularity of individual words.
- Cache replacement policy = LRU

a) (5 points) Show the layout of the cache, clearly showing the fields of the cache (the data and metadata). The sketch should show the number of lines in the cache, the amount of data in each line, and the fields of the metadata in each line.

# CS 2200 Spring 2011 Test 2

Name: \_\_\_\_\_ Prism ID: \_\_\_\_\_  
GTID#: 9 \_\_\_\_\_

b) (5 points) Show how the CPU interprets the memory address (i.e., which bits are used as offset, index, and tag).

c) (5 points) Compute the size of the metadata in **each line of the cache** (remember that this is a 2-way set-associative cache). Show your work to get partial credit.